

PISTIS: An Event-Triggered Real-Time Byzantine-Resilient Protocol Suite

David Kozhaya¹, Jérémie Decouchant, Vincent Rahli, and Paulo Esteves-Verissimo, *Fellow, IEEE*

Abstract—The accelerated digitalisation of society along with technological evolution have extended the geographical span of cyber-physical systems. Two main threats have made the reliable and real-time control of these systems challenging: (i) uncertainty in the communication infrastructure induced by scale, and heterogeneity of the environment and devices; and (ii) targeted attacks maliciously worsening the impact of the above-mentioned communication uncertainties, disrupting the correctness of real-time applications. This article addresses those challenges by showing how to build distributed protocols that provide both real-time with practical performance, and scalability in the presence of network faults and attacks, in probabilistic synchronous environments. We provide a suite of real-time Byzantine protocols, which we prove correct, starting from a reliable broadcast protocol, called *PISTIS*, up to atomic broadcast and consensus. This suite simplifies the construction of powerful distributed and decentralized monitoring and control applications, including state-machine replication. Extensive empirical simulations showcase *PISTIS*'s robustness, latency, and scalability. For example, *PISTIS* can withstand message loss (and delay) rates up to 50 percent in systems with 49 nodes and provides bounded delivery latencies in the order of a few milliseconds.

Index Terms—Real-time distributed systems, probabilistic losses, consensus, atomic broadcast, Byzantine resilience, intrusion tolerance

1 INTRODUCTION

THE accelerated digitalisation of society has significantly shifted the way that physical infrastructures—including large continuous process plants, manufacturing shop-floors, power grid installations, and even ecosystems of connected cars—are operated nowadays. Technological evolution has made it possible to orchestrate a higher and finer degree of automation, through the proliferation of multiple sensing, computing, and communication devices that monitor and control such infrastructures. These monitoring and control devices are distributed by nature of the geographical separation of the physical processes they are concerned with. The overall systems, i.e., the physical infrastructures with their monitoring and control apparatus, are generally known as *cyber-physical systems* (CPS) [1]. However, transposing the monitoring and control functionality normally available in classical, real-time (i.e., adhering to given time bounds) and embedded systems, to the distributed CPS scenarios mentioned above, is a very challenging task, due to two main reasons.

First, the scale of the systems as well as the heterogeneity of devices (sensors, actuators and gateways), induce uncertainty in the communication infrastructure interconnecting them, itself often diverse too, e.g., Bluetooth, Wireless IEEE 802.11, or Fiber [2], [3], [4], [5]. These communication uncertainties become evident [3], [4], [5], namely in the form of link faults and message delays, which hamper the necessary reliability and synchronism needed to realize real-time operations, be it when fetching monitoring data or when pushing decisions to controllers.

Second, security vulnerabilities of many integrated devices, as well as the criticality of the managed physical structures, increase the likelihood of targeted attacks [6], [7]. Such attacks can aim to inflict inconsistencies across system components or to disrupt the timeliness and correctness of real-time applications. The consequences of such attacks can range from loss of availability to severe physical damage [8].

This paper addresses the challenges above, which render traditional approaches for building real-time communications, ineffective in wide-scale, uncertain, and vulnerable settings. We investigate, in particular, how to build large-scale distributed protocols that can provide real-time communication guarantees and can tolerate network faults and attacks, in probabilistic synchronous environments. These protocols simplify the construction of powerful distributed monitoring and control applications, including state-machine replication for fault tolerance. To our knowledge, literature, with the exception of [9], [10], has targeted achieving either real-time guarantees or Byzantine-resilience with network uncertainties, but not both.

To bridge this gap, we present a protocol suite of real-time Byzantine protocols, providing several message delivery

- David Kozhaya is with ABB Research, 5405 Baden, Switzerland.
E-mail: david.kozhaya@ch.abb.com.
- Jérémie Decouchant is with TU Delft, 2628 CD, Delft, Netherlands.
E-mail: j.decouchant@tudelft.nl.
- Vincent Rahli is with the University of Birmingham, B15 2TT Birmingham, U.K. E-mail: vincent.rahli@gmail.com.
- Paulo Esteves-Verissimo is with the King Abdullah University of Science and Technology - RC3, Thuwal 23955, Saudi Arabia.
E-mail: paulo.verissimo@kaust.edu.sa.

Manuscript received 2 Mar. 2020; revised 17 Nov. 2020; accepted 15 Jan. 2021. Date of publication 4 Feb. 2021; date of current version 25 Mar. 2021.
(Corresponding author: David Kozhaya.)
Recommended for acceptance by L. Y. Chen.
Digital Object Identifier no. 10.1109/TPDS.2021.3056718

semantics, from reliable broadcast (*PISTIS*¹), through consensus (*PISTIS-CS*), to atomic broadcast (*PISTIS-AT*). *PISTIS* is capable of: (i) delivering real-time practical performance (i.e., correct nodes provide guarantees within given time bounds) in the presence of aggressive faults and attacks (i.e., one third of the nodes being Byzantine, and high message loss rates); and (ii) scaling with increasing system size.

The main idea underlying *PISTIS* is an event-triggered signature based approach to constantly monitor the network connectivity among processes. Connectivity is measured thanks to the broadcast messages: processes embed signed monitoring information within the messages of the broadcast protocol and exclude themselves from the protocol when they are a threat to timeliness. Hence, *PISTIS* does not modularly build on membership/failure detector oracles (like in traditional distributed computing) but rather directly incorporates such functionalities within. In fact, modularity in this sense was proven to be impossible for algorithms implementing *PISTIS*-like guarantees [10]. In order to mask network uncertainties in a scalable manner, *PISTIS* uses a temporal and spatial gossip-style message diffusion with fast signature verification schemes.

We empirically show that *PISTIS* is robust. For example *PISTIS* can tolerate message loss rates of up to 40, 50, 60, and 70 percent in systems with 25, 49, 73, and 300 nodes respectively: *PISTIS* has a negligible probability of being unavailable under such losses. We also show that *PISTIS* can meet the strict timing constraints of a large class of typical CPS applications, mainly in Supervisory Control And Data Acquisition (SCADA) and Internet of Things (IoT) areas, e.g., (1) fast automatic interactions ($\leq 20\text{ms}$) for systems with up to 200 nodes, (2) power systems and substation automation applications ($\leq 100\text{ms}$) for systems with up to 1,000 nodes, and (3) slow speed auto-control functions ($\leq 500\text{ms}$), continuous control applications ($\leq 1\text{s}$) as well as operator commands of SCADA applications ($\leq 2\text{s}$) for systems with 1,000 nodes or more. Such SCADA and IoT applications could include up to hundreds of devices where reliable and timely communication is required.

By using *PISTIS* as the baseline real-time Byzantine reliable broadcast protocol, we prove that (and show how) higher-level real-time Byzantine resilient abstractions can be modularly implemented, namely, consensus and atomic broadcast. Interestingly, we prove that this can be realized with negligible effort: (1) we exhibit classes of algorithms which are amenable to real-time operations by re-using existing synchronous algorithms from the literature; and (2) we rely on *PISTIS*, which addresses and tolerates the most relevant problems posed by the communication environment, including the impossibility of modularly handling membership/failure detection [10].

In short, our contributions are:

- The *PISTIS* protocol suite, which is to the best of our knowledge the first generic and modular protocol suite that provides message delivery guarantees for protocols ranging from Byzantine reliable broadcast to Byzantine atomic broadcast. *PISTIS* itself is an

event-triggered real-time Byzantine reliable broadcast algorithm that has higher scalability and faster message delivery than conventional time-triggered real-time algorithms, in the presence of randomized and unbounded network disruptions. Building on top of *PISTIS*, we present classes of algorithms, *PISTIS-CS* and *PISTIS-AT*, that implement real-time Byzantine consensus and atomic broadcast, respectively.

- Correctness proofs of the *PISTIS* protocol suite. We provide the main proof results in this paper (exhaustive proofs are deferred to Appx. B, which can be found on the Computer Society Digital Library at <http://doi.ieeecomputersociety.org/10.1109/TPDS.2021.3056718>).
- Extensive empirical simulations using Omnet++ [11] that showcase *PISTIS*'s robustness, latency, and scalability.

Roadmap. The rest of the paper is organized as follows. Section 2 discusses related work. Section 3 details our system model. Section 4 recalls the properties of a real-time Byzantine reliable broadcast, and presents our algorithm, *PISTIS*, in details. Section 5 shows and proves how real-time Byzantine atomic broadcast and consensus can be realized on top of *PISTIS*'s guarantees using classes of existing algorithms. Section 6 evaluates the performance and reliability of *PISTIS*. Finally, Section 7 concludes the paper. For space limitations, proofs and additional material are deferred to Appendices, available online.

2 RELATED WORK

Reliable broadcast is a standard abstraction to ensure that the (correct) nodes of a distributed system agree on the delivery of messages even in the presence of faulty nodes. Byzantine reliable broadcast in particular guarantees that (correct) nodes agree even in the presence of arbitrary faults. It is a key building block of reliable distributed systems such as Byzantine Fault-Tolerant State Machine Replication protocols, which are nowadays primarily used in blockchain systems. Pioneered by the work of Dolev [12] and Bracha [13], many protocols have been proposed since then that are intended to work in various environments. The focus of our paper is on novel Byzantine broadcast primitives and protocols that achieve timeliness guarantees.

This paper has evolved from, and improved over, a research line paved by [9], [10], [14] on timing aspects of reliable broadcast and Byzantine algorithms. Besides these works, the literature on broadcast primitives, to the best of our knowledge, either does not take into account timeliness and maliciousness or addresses them separately.

Cristian *et al.* [9] assumed that all correct processes remain synchronously connected, regardless of process and network failures. This strong network assumption is too optimistic, both in terms of scale and timing behaviour, which in practice leads to poor performance (latency of approximately 2.4 seconds with 25 processes—see Table 1 in Section 6.5 for more details). Moreover, Cristian *et al.*'s system model does not allow processes that malfunction (e.g., by violating timing assumptions) to know that they are treated as faulty by the model. Our algorithm, in

1. *PISTIS* was a Greek goddess who represented the personified spirit (daimona) of trust, honesty and good faith.
Authorized licensed use limited to: Universidad Nacional de Colombia (UNAL). Downloaded on March 31, 2024 at 18:21:30 UTC from IEEE Xplore. Restrictions apply.

TABLE 1
Worst Case Latencies

	[9]	[10]	PISTIS
$N = 25, f = 8$	2,400 ms	26 ms	25.6 ms
$N = 50, f = 16$	8,640 ms	70 ms	27 ms
$N = 100, f = 33$	34,650 ms	150 ms	30 ms

comparison, provides latencies in the range of few milliseconds and our model makes processes aware of their untimeliness.

Verissimo *et al.* [14] addressed the timeliness problem by *weak-fail-silence*: despite the capability of the transmission medium to deliver messages reliably and in real-time, the protocol should not be agnostic of potential timing or omission faults (even if sporadic). The bounded omissions assumption (pre-defined maximum number of omissions) of [14] could not be taken as is, if we were to tolerate higher and more uncertain faults (as we consider in this paper): it could easily lead to system unavailability in faulty periods. Hence we operate with much higher uncertainty levels (faults and attacks).

Kozhaya *et al.* [10] devised a Byzantine-resilient algorithm that provides an upper bound on the delivery latency of messages. This algorithm is time-triggered and relies on an all-to-all communication that limits the algorithm's scalability. Our work improves over [10] on several points: (i) we reduce the delivery latency (few milliseconds as shown in Figs. 7 and 8 compared to a few hundred as shown in [10, Fig. 8]—see also Table 1 for a comparison of worst case latencies) by adopting an event-triggered approach instead of a round-based one; (ii) we improve the system's scalability (at least 5 times less bandwidth consumption) by adopting a gossip-based dissemination instead of an all-to-all communication; and (iii) we show how real-time broadcast primitives can be modularly used to build real-time Byzantine-resilient high-level abstractions like consensus and atomic broadcast.

Guerraoui *et al.* [15] designed a scalable reliable broadcast abstraction that can also be used in a probabilistic setting where each of its properties can be violated with low probability. They achieve a scalable solution by relying on stochastic samples instead of quorums, where samples can be much smaller than quorums. As opposed to this work, our goal is to design a deterministic abstraction where the property are never violated: the real-time Byzantine-resilient reliable broadcast primitive discussed in Section 4 is deterministic because late processes become passive, and therefore count as being faulty.

In [16], [17], the authors present a Byzantine fault-tolerant SCADA system that relies on the Prime [18], [19] Byzantine Fault Tolerant State Machine Replication [20], [21] (BFT-SMR) protocol protocol to ensure both safety and latency guarantees. As opposed to PISTIS, Prime relies on an asynchronous primary-based BFT-SMR. As opposed to Prime, PISTIS-CS and PISTIS-AT algorithms are designed modularly from a timely reliable broadcast primitive; and PISTIS allows slow connections between any processes in a probabilistic synchronous environment, while Prime relies on the existence of a “stable” timely set of processes.

3 SYSTEM AND THREAT MODEL

3.1 System Model

Processes. We consider a distributed system consisting of a set $\Pi = \{p_0, p_1, \dots, p_{N-1}\}$ of $N > 1$ processes. We assume that processes are uniquely identifiable and can use digital signatures to verify the authenticity of messages and enforce their integrity. We denote by $\sigma_i(v)$ the signature of value v by process p_i . We often write σ_i , when the payload is clear from the context. Processes are synchronous, i.e., the delay for performing a local step has a fixed known bound (note that this does not apply to faulty processes—see below).

Clocks. Processes have access to local clocks with a bounded and negligible rate drift to real time. These clocks do not need to be synchronized.

Communication. Every pair of processes is connected by two logical uni-directional links, e.g., p_i and p_j are connected by links l_{ij} and l_{ji} . Links can abstract a physical bus or a dedicated network link. We assume a *probabilistic synchronous communication model*. This means that in any transmission attempt to send a message over on link l_{ij} (with $i \neq j$) at some time t , there is a probability $P_{ij}(t)$ that the message reaches its destination and within a maximum delay d (known to the processes). d is the upper time bound on non-lossy message delivery and $\epsilon_1 < 1 - P_{ij}(t) < \epsilon_2 \ll 1$ where ϵ_1 and ϵ_2 are small strictly positive values. Such violations exist in networks, as arguably all communication is prone to unpredictable disturbances, e.g., bandwidth limitation, bad channel quality, interference, collisions, and stack overflows [4]. Our probabilistic synchronous communication has been shown to be weaker, in some sense [22], than partial synchrony [23]. We further discuss and compare our model to existing traditional ones in Appx. A, available online. We do not model correlated losses explicitly, as previous works like [10] have shown that such bursts can be mitigated and we leave it up to the applications to define how to deal with late messages (i.e., violating the d delay assumption).

3.2 Threat Model

Processes. We assume that some processes can exhibit arbitrary, a.k.a. *Byzantine*, behavior. Byzantine nodes can abstract processes that have been compromised by attackers, or are executing the algorithm incorrectly, e.g., as a result of some fault (software or hardware). A Byzantine process can behave arbitrarily, e.g., it may crash, fail to send or receive messages, delay messages, send arbitrary messages, etc.

We assume that at most $f = \lfloor \frac{N-1}{3} \rfloor$ processes can be Byzantine. This formula was proved to be an upper bound for solving many forms of agreement in a variety of models such as in non-synchronous models [24], [25].

We allow nodes to become *passive* in case they fail to execute in a timely fashion. As explained in Section 4.3, passive nodes stop executing key events to guarantee timeliness. A process that exhibits a Byzantine behavior or that enters the passive mode (see Section 4.3) is termed *faulty*. Otherwise, the process is said to be *correct*. Note that passive nodes are considered faulty (at least) during the time they are passive, but are not counted against the f Byzantine faults. Therefore, more than f nodes could be faulty in a system over the

full lifespan of a system (up to f nodes could be Byzantine, and up to N processes could be momentarily passive).

Clocks. The bounded and negligible rate drift assumption in Section 3.1 has to hold only on a per protocol execution basis, easily met by current technology (such as techniques relying on GPS [26] or trusted components [27]). Hence the clock of a correct process always behaves as described in Section 3.1.

Communication. We assume that Byzantine processes or network adversaries cannot modify the content of messages sent on a link connecting correct processes (implemented by authentication through unforgeable signatures [28]).

4 REAL-TIME BYZANTINE RELIABLE BROADCAST

We now present our solution to guarantee that correct nodes reliably deliver broadcast messages in a timely fashion, despite Byzantine nodes, and communication disruptions. Section 4.1 recalls the properties of the real-time Byzantine-resilient reliable broadcast (RTBRB) primitive [10]. Then, Section 4.2 presents a high-level overview of the PISTIS event-triggered algorithm, which implements the RTBRB primitive, while Section 4.3 provides a detailed presentation of PISTIS. Finally, Section 4.5 explains how passive nodes can recover and become active again to ensure the liveness of the system.

4.1 Real-Time Byzantine Reliable Broadcast Abstraction

Definition 1 (RTBRB). *The real-time Byzantine reliable broadcast primitive guarantees the following properties [10], assuming every message is uniquely identified (e.g., using the pair of a sequence number and a process id—the broadcaster's id).² In this abstraction, a process broadcasts a message by invoking RTBRB-broadcast(). Similarly, a process delivers a message by invoking RTBRB-deliver().*

- **RTBRB-Validity:** *If a correct process p broadcasts m , then some correct process eventually delivers m .*
- **RTBRB-No duplication:** *No correct process delivers message m more than once.*
- **RTBRB-Integrity:** *If some correct process delivers a message m with sender p_i and process p_i is correct, then m was previously broadcast by p_i .*
- **RTBRB-Agreement:** *If some correct process delivers m , then every correct process eventually delivers m .*
- **RTBRB-Timeliness:** *There exists a known Δ_R such that if a correct process broadcasts m at real-time t , no correct process delivers m after real time $t + \Delta_R$.*

It is important to note that the above abstraction does not enforce ordering on the delivery of messages sent. We elaborate more on that and how to achieve order in Section 5. Note also that in a system consisting of correct and faulty nodes, these properties ensure that correct nodes deliver broadcast messages within a bounded delay, while no such guarantee is (and can be) provided about faulty nodes.

2. RTBRB's properties are equivalent to the ones of the Byzantine reliable broadcast abstraction defined in [29, Module 3.12,p.117], excluding *Timeliness*.

4.2 Overview of PISTIS

This section presents a high-level description of *PISTIS*. For simplicity, we assume the total number of processes to be $N = 3f + 1$, in which case a Byzantine quorum has a size of $2f + 1$. *PISTIS* guarantees RTBRB properties deterministically despite the probabilistic lossy network. However, this comes at the price of *PISTIS* triggering an entire system fail-safe (shutdown) and a reinitialization of system state when violating RTBRB-Timeliness is inevitable. We show later in Section 6 that the probability of *PISTIS* causing such system fail-safe (and hence violating an RTBRB property if fail-safe was not triggered) is negligible.

System Awareness. Given that broadcasts can be invoked at unknown times, there might exist a correct process in $\Pi \setminus \{p_i\}$ that is unaware of p_i 's broadcast for an unbounded amount of time after it was issued, since all links can lose an unbounded number of messages. The occurrence of such scenarios may hinder the system's ability of delivering real-time guarantees. To this end, we require that every process p_j constantly exchanges messages with the rest of the system. This regular message exchange aims at capturing how well p_j is connected to other processes, and hence to what extent p_j is up-to-date with what is going on in the system (and to what extent the system knows about p_j 's state). We achieve this constant periodic message exchange via a function, which we call *proof-of-connectivity*.³ It requires each process to diffuse heartbeats to the rest of the system in overlapping rounds: a new round is started every d time units, and each round is of a fixed duration T , where $d < T$. Section 6 shows that $T = 8d$ is a reasonably good value, while Section 4.4 highlights the need for overlapping rounds. A round consists in repeatedly (every d units of time) diffusing a signed heartbeat message to X other processes. X stands for the number of processes to which a process sends a message in a communication step. The value of X is fixed at deployment time (i.e., does not change over the execution of a system) and can range between 0 and $N - 1$. It is used to avoid network congestions by enforcing that processes selectively send their messages to an arbitrary subset of the system. Each round consists then in repeatedly sending $\lceil \frac{T}{d} \rceil$ times a message, each time to X other nodes. Note that even though the value of X is fixed, in any given round the set of X processes to which the message is sent in every repetition can change such that the union of processes to which the message is sent in all $\lceil \frac{T}{d} \rceil$ repetitions in that round covers all processes in the system. This is possible when $N \leq X \times \lceil \frac{T}{d} \rceil$, which we always guarantee in practice. Heartbeat messages are uniquely identified by sequence numbers, which are incremented prior to each round. On receipt of a heartbeat message, a correct process appends its own signature to it as well as all other seen signatures relative to that heartbeat; and sends it to X other processes. At the end of each round, if a process does not receive at least $2f + 1$ signatures (including its own) on its own heartbeat, it enters the passive mode.

Fig. 1 provides an example of a run of the proof-of-connectivity protocol, depicted as a message sequence diagram,

3. Periodic message exchange (heartbeats) has been used to discover the network state in many monitoring algorithms [30], [31]




Fig. 1. Example of a proof-of-connectivity run, where $X = 2f + 1$, and where 2 repetitions allow covering all nodes.

in a system composed of 4 processes. This figure depicts part of the three first rounds of proof-of-connectivity initiated by p_0 (we only show the messages sent by p_0 to avoid cluttering the picture), namely PoC_0 in blue, PoC_1 in orange, and PoC_2 in purple. In addition, in that case, each proof of connectivity round is of length $\mathbb{T} = 6d$. Therefore, the blue PoC_0 heartbeats are sent 6 times between d_0 and d_5 , the orange PoC_1 heartbeats are sent 6 times between d_1 and d_6 , and the purple PoC_2 heartbeats are sent 6 times between d_2 and d_7 . If by the end of PoC_0 , p_0 has not received $2f$ replies to its heartbeats, it will become passive.

Diffusing Broadcasts. PISTIS relies on two types of messages (Echo and Deliver messages) to ensure that broadcast values are delivered in a timely fashion. Processes exchange Echo messages either to start broadcasting new values, or in response to received Echo messages. Echo messages help processes gather a valid quorum (a Byzantine write quorum [32] of size $2f + 1$) of signatures on a single value v relative to a broadcast instance. A broadcast instance is identified by the id of the process broadcasting v and a sequence number. Echo messages help prevent system inconsistencies when malicious nodes send different values with the same sequence number (same broadcast instance) to different recipients. However, additional messages, namely Deliver messages, are needed to help achieve delivery within a bounded time after the broadcast.

When a process p_i receives a value v through an Echo message, it appends its signature to the message as well as all other signatures it has received relative to v ; and sends it to X other processes. In addition, when p_i receives a value for the first time, it triggers a local timer of duration \mathbb{T} . Upon receiving a value signed by more than $2f$ processes, a process delivers that value. However, a process that does not receive more than $2f$ signatures on time (i.e., before the timer expires) enters the passive mode. In case multiple values are heard relative to a single process and sequence number (equivocation), then the first heard value is the one to be echoed. Note that processes continue executing the proof-of-connectivity function during the *echo* and *deliver* phases however by piggybacking heartbeats to echo/deliver messages.

As opposed to Echo messages that are diffused (i.e., retransmitted temporally and sporadically) for a duration \mathbb{T} , Deliver messages are diffused for $2\mathbb{T}$. This is needed to ensure that if some correct processes start diffusing a message between some time t and $t + \mathbb{T}$, possibly at different times, then there must be a \mathbb{T} -long period of time where all of them are diffusing the message (see Lemma 4 in Appx. B for more details), available online. Given a large enough collection of such processes ($f + 1$ correct processes), this




Fig. 2. Example of a PISTIS run where $X = 2f + 1$, and where 2 repetitions allow covering all nodes.

allows other processes to learn about delivered values in a timely fashion.

Algorithm 1. *proof-of-Connectivity*(\mathbb{T}) @ Process p_i

```

1:  $seq = [0]^n$ ; // stores smallest valid sequence number per process.
2:  $sq = 0$ ; // local sequence number.
3:  $\mathcal{R}_{HB} = [0]^n$ ; // stores signatures on last  $\lceil \frac{\mathbb{T}}{d} \rceil$  heartbeats of processes.
4:
5: upon event initialization()  $\vee$  check-connectivity() do
6:   trigger Timeout(msg,  $\mathbb{T}$ );
7:   Execute h-diffuse( $\langle p_i, sq \rangle, \{\sigma_i\}$ );
8:    $\mathcal{R}_{HB}[p_i].add(\langle p_i, sq \rangle, \{\sigma_i\})$ ;  $sleep(d)$ ;  $sq++$ ;
9:   if  $sq - seq[p_i] > \lceil \frac{\mathbb{T}}{d} \rceil$  then  $seq[p_i]++$ ;
10:  end if
11:  trigger check-connectivity();
12:
13: upon event Expired-Timer( $\langle p_i, sq' \rangle, timeout$ ) do
14:   if  $|\mathcal{R}_{HB}[p_i].getsig(sq')| \leq 2f$  then
15:     // gets signatures on message with sequence number  $sq'$ 
16:     Initiate passive mode;
17:   else  $\mathcal{R}_{HB}[p_i].remove(sq')$ ; // remove entry with seq. num.  $sq'$ 
18:   end if
19:
20: upon event receive HB( $\langle p_j, sq' \rangle, \Sigma$ ) do
21:   if ( $sq' \geq seq[p_j]$ ) then
22:      $\mathcal{R}_{HB}[p_j].setsig(sq', \mathcal{R}_{HB}[p_j].getsig(sq') \cup \Sigma \cup \{\sigma_i\})$ ;
23:     if  $j \neq i \wedge sq' \neq seq[p_j]$  then
24:       Execute h-diffuse( $\langle p_j, sq' \rangle, \mathcal{R}_{HB}[p_j].getsig(sq')$ );
25:     end if
26:   end if
27:   if  $sq' > (seq[p_j] + \lceil \frac{\mathbb{T}}{d} \rceil) \wedge j \neq i$  then
28:      $seq[p_j] = sq' - \lceil \frac{\mathbb{T}}{d} \rceil$ ;
29:      $\mathcal{R}_{HB}[p_j].remove(sq'')$ ,  $\forall sq'' < seq[p_j]$ ;
30:   end if
31:
32: Function h-diffuse(msg,  $\Sigma$ )
33:   for (int  $i = 0$ ;  $i \leq \lceil \frac{\mathbb{T}}{d} \rceil$ ;  $i++$ ) do
34:     send HB(msg,  $\Sigma$ ) to  $X$  other processes;
35:      $sleep(d)$ ;
36:   end for
37:

```

Fig. 2 provides an example of a run of PISTIS, depicted as a message sequence diagram. The system is composed of 4 processes. This figure depicts part of the echo (in blue) and deliver (in orange) phases of one broadcast initiated by p_0 (for the purpose of this illustration, only the messages sent by p_0 are shown). The purple “broadcast” and “deliver” tags indicate the times at which p_0 initiated its broadcast, and delivered it. In this example, the echo phase is initially meant to last for a duration of $\mathbb{T} = 6d$. However, it happens here that p_0 received $2f$ echo messages for its broadcast by

$3d + k$, where $0 < k < d$, which is why d_3 is shorter than the other intervals. Therefore, p_0 stops its echo phase and starts its deliver phase at $3d + k$. As mentioned above, the deliver phase lasts for $2\mathbb{T}$. If p_0 has not received $2f$ deliver messages in return by the end of that deliver phase, then it becomes passive.

Algorithm 2. PISTIS @ Process p_i

```

1: Execute proof-of-connectivity( $\mathbb{T}$ );
2:
3: upon event RTBRB-broadcast( $p_i, sq, v$ ) do
4:   Execute proof-of-connectivity in piggyback mode;
5:   Initialize  $\mathcal{R}_{echo}(p_i, sq, v) = \{\sigma_i\}$ ;
6:   Execute b-diffuse( $\langle p_i, sq, v \rangle, \mathbb{T}, echo$ );
7:
8: upon event receive Echo( $\langle p_j, sq, v \rangle, \Sigma$ ) do
9:   if  $\nexists \mathcal{R}_{echo}(p_j, sq, \dots)$  then
10:    Initialize  $\mathcal{R}_{echo}(p_j, sq, v) = \{\sigma_i\} \cup \Sigma$ ;
11:    Execute proof-of-connectivity in piggyback mode;
12:    if  $|\mathcal{R}_{echo}(p_j, sq, v)| \leq 2f$  then
13:      Execute b-diffuse( $\langle p_j, sq, v \rangle, \mathbb{T}, echo$ );
14:    else Execute deliver-msg( $p_j, sq, v, \mathcal{R}_{echo}(p_j, sq, v)$ );
15:    end if
16:  else if  $\exists \mathcal{R}_{echo}(p_j, sq, v)$  then
17:     $\mathcal{R}_{echo}(p_j, sq, v) = \mathcal{R}_{echo}(p_j, sq, v) \cup \Sigma$ ;
18:    if  $|\mathcal{R}_{echo}(p_j, sq, v)| > 2f$  (for the first time) then
19:      Execute deliver-msg( $p_j, sq, v, \mathcal{R}_{echo}(p_j, sq, v)$ );
20:    end if
21:  else if  $\exists \mathcal{R}_{echo}(p_j, sq, v' \neq v)$  then
22:    //  $p_j$  has lied about message with  $sq$ 
23:    if  $|\Sigma| > 2f$  then
24:      remove  $\mathcal{R}_{echo}(p_j, sq, v')$ ;
25:       $\mathcal{R}_{echo}(p_j, sq, v) = \Sigma$ ;
26:      Execute deliver-msg( $p_j, sq, v, \Sigma$ );
27:    end if
28:  end if
29:
30: upon event receive Deliver( $\langle p_j, sq, v, \Sigma \rangle, \Sigma'$ ) do
31:   if  $\nexists \mathcal{R}_{deliver}(p_j, sq, v)$  then
32:     $\mathcal{R}_{echo}(p_j, sq, v) = \mathcal{R}_{echo}(p_j, sq, v) \cup \Sigma$ ;
33:    Execute deliver-msg( $p_j, sq, v, \Sigma$ );
34:   end if
35:    $\mathcal{R}_{deliver}(p_j, sq, v) = \mathcal{R}_{deliver}(p_j, sq, v) \cup \Sigma'$ ;
36:
37: upon event Expired-Timer( $msg, timeout, mode$ ) do
38:   if  $\exists \mathcal{R}_{mode}(msg) \wedge |\mathcal{R}_{mode}(msg)| \leq 2f$  then
39:     switch mode do
40:       case echo
41:         if no lie is discovered on  $msg$  then
42:           Initiate passive mode;
43:         end if
44:       case deliver
45:         Initiate passive mode;
46:     end if
47:
48: Function b-diffuse( $msg, timeout, mode$ )
49:   trigger Timeout( $msg, timeout, mode$ );
50:   for (int  $i = 0$ ;  $i \leq \lceil \frac{timeout}{d} \rceil$ ;  $i++$ ) do
51:      $\Sigma = \mathcal{R}_{mode}(msg)$ ;
52:     switch mode do
53:       case echo
54:         send Echo( $msg, \Sigma$ ) to  $X$  random processes;

```

```

55:       case deliver
56:         send Deliver( $msg, \Sigma$ ) to  $X$  random processes;
57:       sleep( $d$ );
58:     end for
59:
60: Function deliver - msg $p_i$ ( $p_j, sq, v, \Sigma$ )
61:   if  $\nexists \mathcal{R}_{deliver}(p_j, sq, v)$  then
62:     Execute proof-of-connectivity in piggyback mode;
63:     trigger RTBRB-deliver( $p_j, sq, v$ );
64:     Initialize  $\mathcal{R}_{deliver}(p_j, sq, v) = \{\sigma_i\}$ ;
65:     Stop sending any Echo();
66:   end if
67:   Execute b-diffuse( $\langle p_j, sq, v, \Sigma \rangle, 2\mathbb{T}, deliver$ );
68:

```

4.3 Detailed Presentation of PISTIS

We now discuss PISTIS (Algorithm 2) in more details. Note that all functions presented in Algorithms 1 and 2 are non-blocking. PISTIS's proof of correctness can be found in Appx. B, available online.

Process States. Processes can become passive under certain scenarios by calling “Initiate passive mode”. A passive node stops broadcasting and delivering messages to guarantee timeliness but otherwise keeps on replying to messages to help other processes. Processes that were behaving correctly thus far, are considered faulty when they initiate a passive mode and can notify the application above of this fact. Later in this section, we show how processes in the passive mode can come back to normal operation by calling “Initiate active mode”.

Ensuring Sufficient Connectivity. In PISTIS every process executes the *proof-of-connectivity* Algorithm 1. Namely, a process p_i forms a heartbeat $HB(\langle p_i, sq \rangle, \{\sigma_i\})$, where sq is p_i 's current heartbeat sequence number and σ_i is p_i 's signature on $\langle p_i, sq \rangle$. Process p_i also stores (in array \mathcal{R}_{HB}) for every process (including itself) all signatures it receives on heartbeats with a *valid sequence number*. A valid heartbeat sequence number for some process p_j is a sequence number $\geq seq[p_j]$. Heartbeats with lower sequence numbers are simply ignored. To avoid receiving heartbeats from older rounds, we update $seq[p_j]$ every time a heartbeat with a sequence number over $seq[p_j] + \lceil \frac{\mathbb{T}}{d} \rceil$ is received (lines 27–28). After forming its heartbeat, p_i sets a timeout of duration \mathbb{T} , and sends this heartbeat to $X > f$ random processes $\lceil \frac{\mathbb{T}}{d} \rceil$ times (lines 32–36). Process p_i increments its heartbeat sequence number and repeats this whole procedure every $d < \mathbb{T}$. Upon incrementing its heartbeat sequence number, p_i updates its own valid heartbeat sequence numbers (lines 9–10).

A process p_i receiving $HB(\langle p_j, sq' \rangle, \Sigma)$ ignores this heartbeat if sq' is smaller than the smallest valid heartbeat sequence number known for p_j . Otherwise, p_i updates p_j 's valid heartbeat sequence numbers (lines 27–30) and the list of all seen signatures on these valid heartbeats (line 22). Then, p_i diffuses the heartbeat with the updated list of seen signatures to X random processes (line 24).

When a timer expires, p_i checks $\mathcal{R}_{HB}[p_i]$ for the number of accumulated signatures on its corresponding heartbeat. If that number is $\leq 2f$, p_i enters the passive

mode; otherwise it removes the corresponding entry from $\mathcal{R}_{HB}[p_i]$ (lines 13–19).

Broadcasting a Message. A process p_i that wishes to broadcast a value v , calls $\text{RTBRB-broadcast}(p_i, sq, v)$ from Algorithm 2 (lines 3–7), where sq is a sequence number that uniquely identifies this broadcast instance. Given such an event, p_i produces a signature σ_i for the payload $\langle p_i, sq, v \rangle$. It then triggers a timeout of duration \mathbb{T} and sends an $\text{Echo}(\langle p_i, sq, v \rangle, \{\sigma_i\})$ message $\lceil \frac{\mathbb{T}}{d} \rceil$ times to X other random processes. *Proof-of-connectivity* information from p_i is now piggybacked on these messages, as on all other Echo and Deliver messages.

Sending and Receiving Echoes. When p_i receives an $\text{Echo}(\langle p_j, sq, v \rangle, \Sigma)$, p_i reacts differently depending on whether it is not already echoing for this instance (lines 8–15), already echoing v (lines 16–20), or already echoing a different value (lines 21–27). In all three cases, p_i starts delivering a message (and stops sending echoes) as soon as at least $2f + 1$ distinct signatures have been collected for that message.

Sending and Receiving Deliver Messages. When p_i receives $\text{Deliver}(\langle p_j, sq, v, \Sigma \rangle, \Sigma')$ for the first time (lines 60–67), it delivers $\langle p_j, sq, v, \Sigma' \rangle$, and sends $\text{Deliver}(\langle p_j, sq, v, \Sigma \rangle, \mathcal{R}_{\text{deliver}}(p_j, sq, v))$ using $\text{b-diffuse}()$. In case that deliver message is not the first one received (lines 30–35), p_i aggregates all seen signatures for $\langle p_j, sq, v \rangle$ in $\mathcal{R}_{\text{deliver}}(p_j, sq, v)$ (all functions that use $\mathcal{R}_{\text{deliver}}(p_j, sq, v)$ now use the new updated value).

Process Passive Mode. When a timeout set by process p_i with parameters $(msg, timeout, mode)$ expires, p_i enters the passive mode if the set \mathcal{R}_{mode} has less than $2f + 1$ distinct signatures, for $mode = \text{deliver}$. For $mode = \text{echo}$, p_i enters passive mode if in addition to \mathcal{R}_{mode} not having $2f + 1$ signatures, p_i did not discover a lie for that broadcast instance.

Remark 1. Any message of the form $\text{Echo}(\langle p_j, sq, v \rangle, \Sigma_1)$ or $\text{Deliver}(\langle p_j, sq, v, \Sigma_2 \rangle, \Sigma_3)$ is termed invalid if: (1) Σ_1 contains an incorrect signature, and similarly for Σ_2 and Σ_3 ; or (2) Σ_1 does not contain a signature from p_j , and similarly for Σ_2 ; or (3) Σ_2 has less than $2f + 1$ signatures. Invalid messages are simply discarded.

Remark 2. We assume that processes sign payloads of the form (p_i, sq, v, E) for echo messages and of the form (p_i, sq, v, D) for deliver messages. We use the E and D tags to distinguish echo and deliver payloads, thereby ensuring that an attacker cannot use echo signatures as deliver signatures. Note that echo signatures are sent as part of deliver messages as a proof that a quorum of processes echoed a certain value.


4.4 PISTIS' Properties

As mentioned at the beginning of this section, PISTIS is correct in the sense that it satisfies all five properties of the RTBRB primitive presented in Section 4.1:

Theorem 1 (Correctness of PISTIS). *Under the model presented in Section 3, the PISTIS algorithm presented in Fig. 2 implements the RTBRB primitive.*

A proof of this theorem can be found in Appx. B, available online. Let us point out here that the Δ_R bound of the RTBRB-Timeliness property turns out to be $3T$.

Let us also highlight the crux of this proof here.



As illustrated above, a correct node p_i that broadcasts a message m a time t is guaranteed to start delivering m by $t_d = t + T$. In addition thanks to the $2T$ delivery period, we are also guaranteed that a collection, called B , of $2f + 1$ nodes, will only deliver m for a T -long period that starts before $t_d + T$. PISTIS's proof-of-connectivity (PoC) mechanism then ensures that any other correct node p_j will execute a PoC round during which a correct node $r \in B$ delivers m to p_j , piggybacked to a heartbeat, thereby guaranteeing that p_j delivers m timely.

In particular, overlapping PoC rounds allow for all correct nodes to have a PoC round that coincide with that T -long period (called D here), during which the correct nodes in B deliver m , thereby allowing all correct nodes to deliver m . If PoC rounds were consecutive and not overlapping, a correct node could miss the deliver message (piggybacked with PoC messages) sent during D if it were to receive PoC messages for a round (i.e., sequence number) s sent before D , and for round $s + 1$ sent after D , thereby staying active while not delivering.

4.5 Byzantine-Resilient Recovery

If process p_i detects that it is executing under bad network conditions, it enters the passive mode and signals the upper application. As a result, p_i stops broadcasting and delivering broadcast messages (by not executing line 3 and line 63) to avoid violating RTBRB-Timelines. However, p_i continues participating in the dissemination of the broadcast and proof-of-connectivity messages to avoid having too many nodes not collecting enough messages and hence becoming passive.

Once the network conditions are acceptable again, p_i can recover and resume delivering broadcast messages. More precisely, a process p_i that enters passive mode at time t can operate normally again if the interval $[t, t + \Delta_R]$ is free of any passive mode initiations. This Δ_R duration ensures that the messages delivered by a recovered process p_i do not violate any RTBRB properties. After a delay Δ_R , nodes will resume their full participation in the protocol, and either deliver messages or stay on hold.

Note that in case of multiple broadcast instances, passive nodes that become active again should learn the latest sequence number of broadcasts for other nodes. Otherwise Byzantine nodes can exploit this to hinder the liveness of the system.

Remark 3. Given that processes can now shift between passive and active modes, we specify our notion of correct

processes as follows. A system run is modeled by a trace of events happening during that run. An event has a time-stamp and a node associated with it. Moreover, an event can either be a correct event or a Byzantine event. Given an algorithm A , a process p is deemed correct w.r.t. A and a trace τ , if: (1) it follows its specification from e_1 , the first correct A -related event (i.e., an event of algorithm A) happening in τ , to e_2 , the last correct A -related event happening in τ ; (2) p 's events between e_1 and e_2 must all be correct; (3) p must also have followed its specification since it last started; and (4) p must never have lost its keys (so that no other node can impersonate p when p follows its specification). The results presented below also hold for this definition of correctness, because correct processes are required to be active through the entire broadcast instance.

This recovery mechanism improves the overall resilience of the system. Indeed, having all processes in passive mode can occur if $2f + 1$ nodes are passive, which is now harder to achieve if nodes can recover sufficiently fast enough.

5 BEYOND A RELIABLE BROADCAST

Unlike liveness in asynchronous reliable broadcast, the RTBRB-Timeliness property (a safety property) introduces a scent of physical ordering. This ordering is due to the fact that timeliness stipulates, for each execution, a termination event to occur “at or before” some Δ_R on the time-line. This said, the reader may wonder to what extent does the real-time Byzantine-resilient reliable broadcast (of Section 4.1) help in establishing total order?

The answer to this question lies in examining what happens to multiple broadcasts issued by the same or by different nodes. When multiple broadcasts interleave, e.g., when they are issued within a period shorter than Δ_R (the upper time bound on delivering a message), messages might be delivered to different processes in different orders. The timeliness property of the real-time Byzantine-resilient reliable broadcast only ensures that a message m that is broadcast at time t is delivered at any time in $[t, t + \Delta_R]$. Thus, to ensure total order on all system events, e.g., for implementing *State Machine Replication*, additional abstractions need to be built on top of the real-time Byzantine-resilient reliable broadcast primitive that we have developed so far.

In this section, we investigate how to modularly obtain such an order on system events while still preserving real-time and Byzantine-resilience. We define two build blocks that build on top of RTBRB, namely the RTBC real-time Byzantine consensus abstraction (Definition 2)—a fundamental building block for state machine replication, atomic broadcast and leader election [30]; and the RTBAB real-time atomic broadcast abstraction (Definition 4)—to establish total order on system events. We then provide characterizations of classes of algorithms that implement these abstractions: Theorem 2 provides a characterization of the PISTIS-CS class of algorithms that implement RTBC, while Theorem 3 provides a characterization of the PISTIS-AT class of algorithms that implement RTBAB. Finally, we provided examples of algorithms that belong to these classes (see Examples 1 and 2).

We start with the following assumption that constrains the ways processes can communicate.

Assumption 1. *Correct processes access the network only via the RTBRB primitive, namely using the two operations: RTBRB-broadcast() and RTBRB-deliver().*

From Assumption 1, a correct process p_i that receives a message from an operation other than RTBRB-deliver() simply ignores that message by dropping it.

5.1 Real-Time Byzantine Consensus

Roughly speaking, solving the *Byzantine consensus* problem requires the agreement of distributed processes on a given value, even though some of the processes may fail arbitrarily. Byzantine consensus was first identified by Pease *et al.* [33], and formalized as the *interactive consistency* problem. An algorithm achieves interactive consistency if it allows the non-faulty processes to come to a consistent view of the initial values of all the processes, including the faulty ones. Once interactive consistency has been reached, the non-faulty processes can reach consensus by applying a deterministic averaging or filtering function on the values of their view. We apply the following assumption to reach consensus.

Assumption 2. *Once interactive consistency terminates, every correct process scans the obtained vector and decides on the value that appears at least $2f + 1$ times. If no such value exists, then the process decides \perp , a distinguished element that indicates that no value has been decided.*

Definition 2 (RTBC). *The real-time Byzantine consensus (RTBC) abstraction is expressed by the following properties:⁴*

- **RTBC-Validity:** *If all correct processes propose the same value v , then any correct process that decides, decides v . Otherwise, a correct process may only decide a value that was proposed by some correct process or \perp .*
- **RTBC-Agreement:** *No two correct processes decide differently.*
- **RTBC-Termination:** *Correct processes eventually decide.*
- **RTBC-Timeliness:** *If a correct process p_i proposes a value to consensus at time t , then no correct process decides after $t + \Delta_C$.*

In RTBC a process p_i can propose a value v to consensus by invoking $\text{RTBC-propose}(p_i, \text{inst}, v)$, where inst is a sequence number that uniquely identifies a RTBC instance. Similarly, a process p_i decides on a value v by invoking $\text{RTBC-decide}(p_i, \text{inst}, v)$. In addition $\text{RTBC-init}(\text{inst})$ instantiate a new instance of RTBC with id inst , i.e., for sequence number inst .

Definition 3. *An algorithm is said to be bounded if it only uses a known bounded number of communication rounds.*

Theorem 2 (Characterization of the PISTIS-CS class).

Let PISTIS-CS be the class of bounded (Definition 3) algorithms that implements interactive consistency under Assumptions 1 and 2. Then, PISTIS-CS algorithms also implement RTBC in our model (described in Section 3).

4. The properties of RTBC are the same as the ones of the traditional (strong) Byzantine consensus defined in [23] (see also [29, Module 5.11, p.246]), excluding the *Timeliness* property.

See Appx. C, available online, for a proof of this result.

Example 1 (Examples of PISTIS-CS algorithms). Because the interactive consistency problem has been solved using different algorithms that satisfy Definition 3, our result applies to various existing algorithms, such as [33], [34], [35], [36].

5.2 Real-Time Byzantine-Resilient Atomic Broadcast

Definition 4 (RTBAB). A real-time Byzantine-resilient atomic broadcast (RTBAB) has the same properties as RTBRB (with a different timeliness bound) plus an additional ordering property (therefore, we only present the properties that differ from RTBRB's):

- **RTBAB-Timeliness:** There exists a known Δ_A such that if a correct process broadcasts m at time t , no correct process delivers m after real time $t + \Delta_A$.
- **RTBAB-Total order:** Let m_1 and m_2 be any two messages and suppose that p_i and p_j are any two correct processes that deliver m_1 and m_2 . If p_i delivers m_1 before m_2 , then p_j delivers m_1 before m_2 .

We now define the class of algorithms (called *RoundBased*), through the properties listed below, that modularly implement RTBAB properties. *RoundBased* algorithms make use of a single RTBRB instance and multiple instances of RTBC. We first constrain a *RoundBased* algorithm to start an RTBRB instance within a bounded amount of time for any broadcast call.

Property 1. If a correct process p_i RTBAB-broadcasts a message m at time t , then it also RTBRB-broadcasts m by time $t + \Delta_B$, for some bounded Δ_B .

We then require a *RoundBased* algorithm to start (or end in case this has already been done before) an RTBC instance, within a bounded amount of time, every time the RTBRB instance delivers.

Property 2. If a correct process RTBRB-delivers a message m at time t , such that m 's broadcaster is also correct, then it either RTBC-proposes or RTBC-decides m by $t + \Delta_P$, for some bounded Δ_P .

In addition, the next property constrains the values that can be proposed at each RTBC instance, namely that at most one non- \perp value can be proposed at each instance.

Property 3. Given an RTBC instance $inst$, there exists a value v , such that each correct process either RTBC-propose v or \perp at $inst$.

Next, we require a *RoundBased* algorithm to deliver a RTBC-decided value within a bounded amount of time (Property 4) and to ensure that non-RTBC-decided values are re-proposed in later RTBC rounds (Property 5).

Property 4. If a correct process RTBC-decides a message m at time t , then it also RTBAB-delivers m by time $t + \Delta_D$, for some bounded Δ_D .

Property 5. A correct process p_i that proposes a value v at a given time t , using a given RTBC instance $inst$, and such that

this instance does not decide v , also RTBC-propose v at some instance $inst + k$, where $0 < k$. Moreover, p_i RTBC-proposes v at the smallest instance between $inst + 1$ and $inst + k$ where m is proposed by some process.

Algorithm 3. Example of a PISTIS-AT Algorithm

@Process p_i

```

1: upon event RTBAB-init(rtbab) do
2:    $unordered = []$ ;  $next = [0]^n$ ;  $seq = 0$ ;
3:    $delivered = \emptyset$ ;  $busy = \text{False}$ ;  $inst = 0$ ;
4:
5: upon event RTBAB-broadcast( $p_i, m$ ) do
6:   trigger RTBRB-broadcast( $p_i, seq, m$ );
7:    $seq++$ ;
8:
9: upon event RTBRB-deliver( $p_j, num, m$ ) do
10:  if  $num = next[p_j]$  then
11:     $next[p_j] = next[p_j] + 1$ ;
12:    if  $m \notin delivered$  then
13:       $unordered[p_j] = unordered[p_j].append(\langle p_j, m \rangle)$ ;
14:    end if
15:  else {wait( $\Delta_w$ ); trigger RTBRB-deliver( $p_j, num, m$ );}
16:  end if
17:
18: upon event  $\exists p_j : unordered[p_j] \neq [] \wedge busy = \text{False}$  do
19:    $busy = \text{True}$ ;
20:   trigger RTBC-init( $inst$ );
21:   // initiate a new real-time Byzantine consensus instance
22:   if  $unordered[leader(inst)] \neq []$  then
23:      $m = unordered[leader(inst)].head()$ ;
24:   else { $m = \perp$ ;}
25:   end if
26:   trigger RTBC-propose( $p_i, inst, m$ );
27:
28: upon event RTBC-decide( $p_i, inst', decided$ ) do
29:   if  $inst' = inst$  then
30:     if  $decided \notin delivered \wedge decided \neq \perp$  then
31:        $delivered = delivered \cup \{decided\}$ ;
32:       trigger RTBAB-deliver( $leader(inst), decided$ );
33:     end if
34:      $unordered[leader(inst)].remove(decided)$ ;
35:      $inst++$ ;  $busy = \text{False}$ ;
36:   else {wait( $\Delta_w$ ); trigger RTBC-decide( $p_i, inst', decided$ );}
37:   end if
38:
39: Function leader( $instance$ ) {return( $instance \bmod n$ );}
40:

```

Finally, we require that nodes participate in all successive RTBC instances in a monotonic fashion.

Property 6. Correct processes RTBC-propose exactly one value per RTBC instance; propose values in all RTBC instances (i.e., for all instances $inst \in \mathbb{N}$); in increasing order w.r.t. the instance numbers of the RTBC instances (i.e., if p_i proposes values at times t_1 and t_2 using the RTBC instances $inst_1$ and $inst_2$, respectively, and $t_1 < t_2$, then $inst_1 < inst_2$); and not in parallel (i.e., if p_i proposes a value at time t using an RTBC instance $inst$, and that this RTBC instance has not decided by time $t' > t$, then p_i does not propose any other value between t and t').

Definition 5. Let *RoundBased* be the class of round-based algorithms that satisfy the Properties 1, 2, 3, 4, 5, and 6.

Theorem 3 (Characterization of the PISTIS-AT class).

Let PISTIS-AT be the class of RoundBased algorithms that implement the traditional Byzantine total-order broadcast under Assumption 1. Then, PISTIS-AT algorithms also implement RTBAB in our system (described in Section 3).

To prove Theorem 3, it is sufficient to prove that a RTBAB-broadcasted value m is always RTBAB-delivered within a bounded amount of time. Because of the round-based property, m must be RTBRB-proposed and RTBRB-decided within a bounded amount of time. Consequently there is (within a bounded amount of time) an RTBC instance where “enough” correct nodes RTBC-propose m , so that m gets RTBC-decided upon and RTBAB-delivered within a bounded amount of time. The proof of Theorem 3 is detailed in Appx. D, available online.

We have introduced bounds for each of the operations executing in bounded time, namely Δ_R (Definition 1), Δ_C (Definition 2), Δ_W (Algorithm 3), Δ_B (Property 1), Δ_P (Property 2), Δ_D (Property 4), and Δ_A (Definition 4). Those bounds are not assumed to be related to each other. However, the bound for Δ_A we exhibit in Theorem 3’s proof is a combination of all the other bounds discussed above.

Example 2 (Example of a PISTIS-AT algorithm). Finally, Algorithm 3 provides an example of a PISTIS-AT algorithm that implements RTBAB modularly, which we adapted from [29, Alg.6.2,p.290] to guarantee timeliness.

6 EVALUATION AND COMPARISON

In this section, we evaluate PISTIS’s reliability, latency, and incurred overhead on network bandwidth.

6.1 PISTIS’s Latency Versus Related Systems’ Latency

We begin with a latency comparison between PISTIS and other related works based on the worst case incurred delay. We compute worst case delays from the bounds established for each algorithm (a direct experimental evaluation would not be fair, since not all previous work [9] consider probabilistic synchronous networks). Later sections provide an experimental comparison with RT-ByzCast [10], the system most related to ours. We elaborate in what follows on the computation of the worst case delays. First we refine the definition of d introduced in Section 3.1. Let d_n be the maximum network delay, and d_p be the maximum local processing time, which includes the cryptographic operations overhead, such that d can be decomposed as $d_p + d_n$. Christian *et al.* [9] compute the worst case delay as $10 * (f + 2) * (n - 1) * d_n$ where f is the maximum number of faulty processes, n the total number of processes, and d_n the network delay. In this work, d_p is equal to 10. Kozhaya *et al.* [10] compute the worst-case delay as $3 * R * d$, where R is the number of consecutive synchronous communication rounds the same message gets disseminated (time-triggered re-transmissions). PISTIS’s worst case delay is proved to be $3 * T$. To ensure fairness and consistency with the latency experiments presented below, we set $R = 8$ and $T = 8d$. However, due to PISTIS’s signature management (see, for example, the optimizations described in Section 6.2), PISTIS’s worst case delay can be alternatively computed as $(3 * 8 * d_n) +$

$(2 * N * d_p)$. This is in part due to the fact that in PISTIS nodes avoid re-verifying already verified signatures.

Our results, shown in Table 1, show that PISTIS has the best worst case latencies of all algorithms for $d_n = 1ms$ (as mentioned above, in the first column $d_p = 10$, while in the last two columns d_p is such that $1 < d_p < 10$, and can be derived from the numbers provided in the table).

Two main observations can be made: (1) compared to the other protocols, PISTIS has superior performance due to the fact that PISTIS is event triggered, utilizes fast signature schemes, reduces the number of signatures created and verified, sends fewer messages (which increase individual message failures) and allows processes for fast detection of their tardiness; and (2) PISTIS’s expected performance in practice (see Fig. 7) is significantly better than the worst case delay bound reported in the table.

6.2 Implementation Optimizations

We implemented three optimizations to improve the performance of PISTIS (as described in Section 4.3). (1) If a process p_i knows that some process p_j has already received $2f + 1$ echo signatures for some message m , p_i stops sending echoes related to m to p_j . Every process implements this optimization by maintaining a list, say \mathcal{L} , that contains all the processes from which it has heard $2f + 1$ signatures for a given message. During a broadcast, a process diffuses a message to X processes at random among $\Pi \setminus \mathcal{L}$. Processes do the same for *deliver* messages. (2) Processes do not verify signatures that they have already received. (3) Processes skip messages that only contain signatures that were already received.

6.3 Implementation Configuration and Settings

We implemented PISTIS in C++ on the Omnet++ 5.4.1 network simulator [12]. In order to accurately measure PISTIS’s communication overhead, we configure network links to have a non-limiting 1 Gbps throughput, and a communication latency of either 1 ms or 5 ms. We evaluated PISTIS’s performance using two signature schemes of similar security guarantees, and available in the OpenSSL library [37]: RSA-2048 (i.e., 256 bytes long signatures) and ECDSA with prime256v1 curves (i.e., 71 bytes long signatures). We use broadcast messages of sizes equal to 1 B and 1 KB.

We run our simulations for systems with $N \in \{25, 49, 73, 300\}$ processes in fully connected networks, and for several values of X , which is the number of processes each process forwards a message m to during diffusion. We consider the probability of losing/omitting a message sent at any point in time to be $i/10$, where $0 \leq i \leq 9$.

6.4 PISTIS’s Reliability

To assess PISTIS’s reliability, we evaluate the probability that a correct process enters the passive mode. Such probability is a crucial measure: a process becoming passive may lead the system to shutdown and hence to stop delivering messages. Namely, when $N = 3f + 1$, a single correct process staying passive for long-enough can, in the worst case (when f Byzantine processes are not sending messages), leave $2f$ correct processes, which would not be enough to




Fig. 3. Probability of a correct process becoming passive when $T = 6d$ or $T = 8d$, and $X = f + 1$ (without recovery).

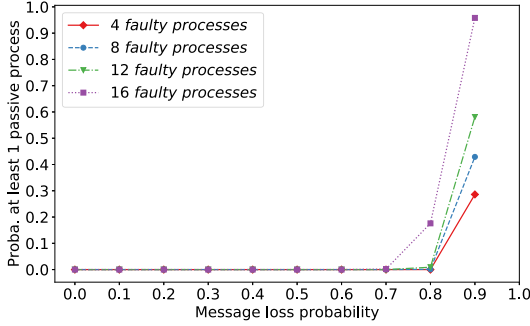


Fig. 4. Probability of a correct process becoming passive in a system of 49 processes (i.e., $f = 16$) using $T = 8d$ and $X = 17$, when 0, 4, 8, 12 or 16 processes are faulty (without recovery).

gather quorums of size $2f + 1$, leading those $2f$ processes to also become passive.

For a given value of N and p , we invoke a broadcast at one of the processes and record any non-Byzantine process that crashed itself during broadcast. We obtain our results by repeating each experiment 10^5 times, and we report the probability that a process crashes itself as

$$(\text{num. of experiments with self-crashed processes})/10^5.$$

We study the impact of several parameters, including T , N , X , f , and p , on PISTIS's reliability, and determine which values should be used to enforce an intended system reliability.

Fig. 3 shows that the system's reliability increases with its size and T 's value for large enough values of T/d . For example, when $T = 8d$, a system with 25 (resp. 49) processes operates with high reliability (i.e., there is a negligible probability that a process becomes passive) under message loss rates reaching up to 40 percent (resp. 50 percent).

Fig. 4 shows that the actual number of Byzantine processes, which varies between 0 and f (the maximum number of tolerable Byzantine nodes), influences the system's resiliency. As expected, with fewer processes being Byzantine, higher message loss rates are tolerated without any process shutdown.

Impact of the Diffusion Fanout. In the results presented so far, processes forward each message to $X = f + 1$ other random processes. We now study the effect of X by measuring PISTIS's reliability when it varies. Fig. 5 shows that increasing X helps increase the overall system reliability. As expected increasing the fanout (value of X) reduces the probability of having a non-Byzantine node becoming passive.

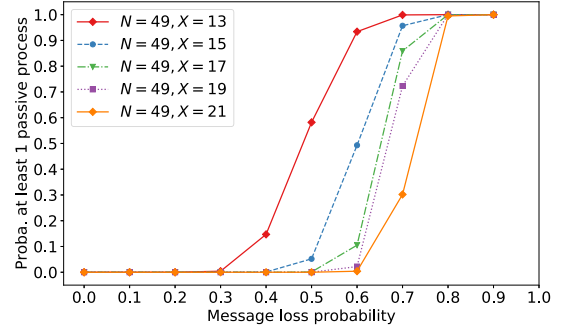


Fig. 5. Probability of a correct process becoming passive in a system of 49 processes using $T = 8d$, and where X varies (without recovery).




Fig. 6. Probability that no Byzantine quorum remains active in systems of 49 or 52 processes, when $T = 8d$, $X = 17$, and $f = 16$ processes are Byzantine.

Recovery. Fig. 6 details the probability that no Byzantine quorum remains active after a broadcast instance when the message loss probability increases. First, one can observe that the recovery mechanisms improve the resiliency of the system. For example, with $N = 49$, PISTIS can tolerate a 70 percent message loss rate without system-wide crashes thanks to the recovery mechanisms, improving over the value of 50 percent obtained without recovery. Second, we show that one can further improve the system's tolerance to message losses by overprovisioning the system. By using three more nodes, i.e., 52 in total, the system can tolerate $f = 16$ Byzantine nodes and now tolerate up to 80 percent of message losses.

6.5 PISTIS Latency and Bandwidth Consumption

Next, we evaluate PISTIS's incurred bandwidth and latency. For these experiments, we average results over 1,000 runs. We use $T = 8d$, since our reliability results show it allows a very large number of message losses to be tolerated. However, we now run our experiments without any message losses to measure the worst case bandwidth consumption. We measure both the protocol latency and bandwidth consumption depending on the value of X that the processes use. We also compare the average latency and bandwidth consumption of PISTIS with that of RT-ByzCast [10]. Note that RT-ByzCast [10] uses ECDSA signatures and all-to-all communication ($X = N$).

Latency. Figs. 7 and 8 detail the latency for a broadcast message to be delivered by all correct processes in systems of size 25, 49, and 73 (i.e., where $f \in \{8, 16, 24\}$): PISTIS delivers with latencies within $[3\text{ms}, 60\text{ms}]$ depending on the network delay d and signature scheme used RSA versus ECDSA. The latency increases when N increases, and




Fig. 7. Average latency with a 1ms link latency with $\mathbb{T} = 8d$ and without message losses. The dotted lines indicate RT-ByzCast's values [10].




Fig. 8. Average latency with a 5ms link latency. The dotted lines indicate RT-ByzCast's values [10].




Fig. 9. Average bandwidth consumption per node and per communication link with a 1ms link latency without message losses. The dotted lines indicate RT-ByzCast's values [10].

decreases when X increases. We draw the following conclusions: (1) PISTIS is slower than RT-ByzCast for $X < f$. For $X \geq f$ PISTIS is on a par with RT-ByzCast until some $X \leq 3f$ ($X \leq 2f$ for systems with up to 400 nodes, see Table 2) after which PISTIS is faster; (2) PISTIS's absolute improvement over RT-ByzCast becomes more significant with increased link delay; (3) When delivering latencies on par with or better than RT-ByzCast, PISTIS can do so with a lower network overhead as presented next (see Figs. 9 and 10).

Network Bandwidth Consumption. We now measure PISTIS's bandwidth overhead per broadcast invocation, using RSA and ECDSA signatures. Figs. 9 and 10 present the bandwidth consumption for 1B payloads with 1 ms and 5 ms link delay, respectively. One can observe that with $X = f + 1$ and when using ECDSA signatures, PISTIS's bandwidth consumption is 3.2 times lower than that of RT-ByzCast. We also observe that when using ECDSA signatures there is a fanout between $2f + 1$ and $3f + 1$ such that below this fanout PISTIS's average




Fig. 10. Average bandwidth consumption per node and per communication link with a 5ms link latency without message losses. The dotted lines indicate RT-ByzCast's values [10].




Fig. 11. Average bandwidth consumption per node and per communication link with a 1ms link latency using either 1B or 1KB messages, without message losses.

bandwidth consumption is lower than RT-ByzCast's, while past that threshold, PISTIS's average bandwidth consumption becomes greater than RT-ByzCast's. This is partly due to the fact that PISTIS being event-based sometimes consumes more bandwidth. However, we see in those figures that PISTIS provides a useful trade-off between latency and bandwidth consumption. Fig. 11 shows as well that the bandwidth consumption increases reasonably when the message payload is increased to 1KB. Besides bandwidth, Fig. 12 (Appx. E), available online, shows that PISTIS also sends less message than RT-ByzCast.

Scalability With the System Size. We also evaluated how PISTIS' latency and bandwidth consumption evolve with larger system sizes, namely up to 1,000 nodes for $X \geq f + 1$ and a 5 ms link latency. Table 2 summarizes the results obtained for $X = f + 1$, $X = 2f + 1$ and $X = N$. Our results show that PISTIS outperforms RT-ByzCast and provides latencies suitable for (1) fast automatic interactions (≤ 20 ms) for systems with up to 200 nodes, (2) power systems and substation automation applications (≤ 100 ms) for systems with up to 1,000 nodes, and (3) slow speed auto-control functions (≤ 500 ms), continuous control applications (≤ 1 s) and operator commands of SCADA applications (≤ 2 s) for systems with 1,000 nodes or more.

7 CONCLUSION

In this paper, we studied how to build large-scale distributed protocols that tolerate network faults and attacks while providing real-time communication. We introduced a suite of proven correct algorithms, starting from a baseline real-time Byzantine

TABLE 2

Pistis Bandwidth Consumption (Mbps) and Broadcast Duration (ms) With Larger Systems ($f = \lfloor N/3 \rfloor$), Where $X_{min} = f + 1$, $X_{mid} = 2f + 1$ and $X_{max} = N$

N	Bdw, X_{min}	Bdw, X_{mid}	Bdw, X_{max}	Bdw [10]	Lat, X_{min}	Lat, X_{mid}	Lat, X_{max}	Lat [10]
25	0.6	1.2	1.7	1.4	21.1	11.0	11.1	20.9
49	1.0	2.2	3.1	2.6	22.3	12.4	12.0	22.0
73	1.5	3.2	4.6	3.9	23.6	13.1	13.2	23.1
200	3.8	8.4	12.5	10.4	31.5	20.7	19.7	29.3
300	5.7	12.5	18.6	15.6	41.2	31.2	27.4	38.0
400	7.6	16.7	25.0	20.9	59.7	43.0	32.0	41.2
500	9.4	20.8	31.1	26.0	85.1	63.0	40.0	51.6
1000	18.7	41.4	62.2	52	296.3	213.1	98.5	116.2

reliable broadcast algorithm, called PISTIS, all the way up to real-time Byzantine atomic broadcast and consensus algorithms. PISTIS is empirically shown to be robust, scalable, and capable of meeting timing deadlines of real CPS applications. PISTIS withstands message loss (and delay) rates up to 50 percent in systems with 49 nodes and provides bounded delivery latencies in the order of a few milliseconds. PISTIS improves over the state-of-the-art in scalability and latency through its event-triggered nature, gossip-based communications, and fast signature verifications. Our work simplifies the construction of powerful distributed and decentralized monitoring and control applications of various CPS domains, including state-machine replication for fault and intrusion tolerance.

ACKNOWLEDGMENTS

The work of Vincent Rahli was supported in part by the National Cyber Security Centre (NCSC) project: Aion: Verification of Critical Components' Timely Behavior in Probabilistic Environments.

REFERENCES

- [1] J. R. Moyne and D. M. Tilbury, "The emergence of industrial control networks for manufacturing control, diagnostics, and safety data," *Proc. IEEE*, vol. 95, no. 1, pp. 29–47, Jan. 2007.
- [2] R. Jacob, M. Zimmerling, P. Huang, J. Beutel, and L. Thiele, "End-to-end real-time guarantees in wireless cyber-physical systems," in *Proc. IEEE Real-Time Syst. Symp.*, 2016, pp. 167–178.
- [3] L. Schenato, B. Sinopoli, M. Franceschetti, K. Poolla, and S. S. Sastry, "Foundations of control and estimation over lossy networks," *Proc. IEEE*, vol. 95, no. 1, pp. 163–187, Jan. 2007.
- [4] D. Dzung, R. Guerraoui, D. Kozhaya, and Y.-A. Pignolet, "To transmit now or not to transmit now," in *Proc. IEEE 34th Symp. Reliable Distrib. Syst.*, 2015, pp. 246–255.
- [5] DLC+VIT4IP, "D1.1 scenarios and requirements specification," Tech. Rep., 2010. [Online]. Available: <http://www.dlc-vit4ip.org/wb/media/Downloads/D1.1-V0.5-20100910-team.pdf>
- [6] M. M. Patel and A. Aggarwal, "Security attacks in wireless sensor networks: A survey," in *Proc. Int. Conf. Intell. Syst. Signal Process.*, 2013, pp. 329–333.
- [7] F. Januário, C. Carvalho, A. Cardoso, and P. Gil, "Security challenges in SCADA systems over wireless sensor and actuator networks," in *Proc. 8th Int. Congress Ultra Modern Telecommun. Control Syst. Workshops*, 2016, pp. 363–368.
- [8] P. Polityuk, O. Vukmanovic, and S. Jewkes, "Ukraine's power outage was a cyber attack: Ukrenerg," 2017. [Online]. Available: <https://www.reuters.com/article/us-ukraine-cyber-attack-energy/ukraines-power-outage-was-a-cyber-attack-ukrenerg-idUSKBN1521BA>
- [9] F. Cristian, H. Aghili, H. R. Strong, and D. Dolev, "Atomic broadcast: From simple message diffusion to byzantine agreement," *Inf. Comput.*, vol. 118, no. 1, pp. 158–179, 1995.
- [10] D. Kozhaya, J. Decouchant, and P. Esteves-Verissimo, "RT-ByzCast: Byzantine-resilient real-time reliable broadcast," *IEEE Trans. Comput.*, vol. 68, no. 3, pp. 440–454, Mar. 2019.
- [11] OMNeT++, Accessed: Mar. 12, 2021. [Online]. Available: <https://omnetpp.org>
- [12] D. Dolev, "Unanimity in an unknown and unreliable environment," in *Proc. 22nd Annu. Symp. Found. Comput. Sci.*, 1981, pp. 159–168. [Online]. Available: <https://doi.org/10.1109/SFCS.1981.53>
- [13] G. Bracha, "Asynchronous byzantine agreement protocols," *Inf. Comput.*, vol. 75, no. 2, pp. 130–143, 1987.
- [14] P. Verissimo, L. Rodrigues, and M. Baptista, "AMP: A highly parallel atomic multicast protocol," in *Proc. Symp. Commun. Architectures Protocols*, 1989, pp. 83–93.
- [15] R. Guerraoui, P. Kuznetsov, M. Monti, M. Pavlovic, and D. Seredinschi, "Scalable byzantine reliable broadcast," in *Proc. Int. Symp. Distributed Comput.*, 2019, pp. 22:1–22:16. [Online]. Available: <https://doi.org/10.4230/LIPIcs.DISC.2019.22>
- [16] A. Babay et al., "Deploying intrusion-tolerant SCADA for the power grid," in *Proc. 49th Annu. IEEE/IFIP Int. Conf. Dependable Syst. Netw.*, 2019, pp. 328–335. [Online]. Available: <https://doi.org/10.1109/DSN.2019.00043>
- [17] A. Babay, T. Tantilillo, T. Aron, M. Platania, and Y. Amir, "Network-attack-resilient intrusion-tolerant SCADA for the power grid," in *Proc. 48th Annu. IEEE/IFIP Int. Conf. Dependable Syst. Netw.*, 2018, pp. 255–266. [Online]. Available: <https://doi.org/10.1109/DSN.2018.00036>
- [18] Y. Amir, B. A. Coan, J. Kirsch, and J. Lane, "Byzantine replication under attack," in *Proc. IEEE Int. Conf. Dependable Syst. Netw. FTCS and DCC*, 2008, pp. 197–206. [Online]. Available: <https://doi.org/10.1109/DSN.2008.4630088>
- [19] Y. Amir, B. A. Coan, J. Kirsch, and J. Lane, "Prime: Byzantine replication under attack," *IEEE Trans. Dependable Secure Comput.*, vol. 8, no. 4, pp. 564–577, Jul./Aug. 2011. [Online]. Available: <https://doi.org/10.1109/TDSC.2010.70>
- [20] F. B. Schneider, "Implementing fault-tolerant services using the state machine approach: A tutorial," *ACM Comput. Surv.*, vol. 22, no. 4, pp. 299–319, 1990. [Online]. Available: <https://doi.org/10.1145/98163.98167>
- [21] M. Castro and B. Liskov, "Practical byzantine fault tolerance," in *Proc. 3rd Symp. Operating Syst. Des. Implementation*, 1999, pp. 173–186. [Online]. Available: <https://dl.acm.org/citation.cfm?id=296824>
- [22] D. Dzung, R. Guerraoui, D. Kozhaya, and Y.-A. Pignolet, "Never say never - Probabilistic and temporal failure detectors," in *Proc. IEEE Int. Parallel Distrib. Process. Symp.*, 2016, pp. 679–688.
- [23] D. Dolev, C. Dwork, and L. Stockmeyer, "On the minimal synchronism needed for distributed consensus," *J. ACM*, vol. 34, no. 1, pp. 77–97, 1987.
- [24] D. Dolev, "The byzantine generals strike again," Stanford Univ., Stanford, CA, Tech. Rep. STAN-CS-81-846, 1981.
- [25] M. J. Fischer, N. A. Lynch, and M. Merritt, "Easy impossibility proofs for distributed consensus problems," *Distrib. Comput.*, vol. 1, no. 1, pp. 26–39, 1986. [Online]. Available: <https://doi.org/10.1007/BF01843568>
- [26] S. Viswanathan, R. Tan, and D. K. Y. Yau, "Exploiting power grid for accurate and secure clock synchronization in industrial IoT," in *Proc. IEEE Real-Time Syst. Symp.*, 2016, pp. 146–156.
- [27] P. Verissimo and A. Casimiro, "The timely computing base model and architecture," *IEEE Trans. Comput.*, vol. 51, no. 8, pp. 916–930, Aug. 2002.
- [28] J. H. An, Y. Dodis, and T. Rabin, "On the security of joint signature and encryption," in *Proc. Int. Conf. Theory Appl. Cryptographic Techn.*, 2002, pp. 83–107.
- [29] C. Cachin, R. Guerraoui, and L. Rodrigues, *Introduction to Reliable and Secure Distributed Programming*. Berlin, Germany: Springer-Verlag, 2011.
- [30] M. K. Aguilera, C. Delporte-Gallet, H. Fauconnier, and S. Toueg, "On implementing omega in systems with weak reliability and synchrony assumptions," *Distrib. Comput.*, vol. 21, no. 4, pp. 285–314, 2008. [Online]. Available: <https://doi.org/10.1007/s00446-008-0068-y>
- [31] R. Guerraoui, D. Kozhaya, and Y. A. Pignolet, "Right on time distributed shared memory," in *Proc. IEEE Real-Time Syst. Symp.*, 2016, pp. 315–326.
- [32] D. Malkhi and M. K. Reiter, "Byzantine quorum systems," in *Proc. 29th Annu. ACM Symp. Theory Comput.*, 1997, pp. 569–578. [Online]. Available: <http://doi.acm.org/10.1145/258533.258650>

- [33] M. Pease, R. Shostak, and L. Lamport, "Reaching agreement in the presence of faults," *J. ACM*, vol. 27, no. 2, pp. 228–234, 1980.
- [34] D. Dolev and H. R. Strong, "Authenticated algorithms for byzantine agreement," *SIAM J. Comput.*, vol. 12, no. 4, pp. 656–666, 1983.
- [35] D. Dolev and R. Reischuk, "Bounds on information exchange for byzantine agreement," *J. ACM*, vol. 32, no. 1, pp. 191–204, 1985.
- [36] L. Lamport, R. Shostak, and M. Pease, "The byzantine generals problem," *ACM Trans. Program. Lang. Syst.*, vol. 4/3, pp. 382–401, 1982.
- [37] OpenSSL. Accessed: Mar. 12, 2021. [Online]. Available: <https://www.openssl.org/>



David Kozhaya received the PhD degree in computer science from EPFL, Switzerland, in 2016, where he was granted a fellowship from the doctoral school. He is a senior scientist at ABB Research, Switzerland. His primary research interests include reliable distributed computing, real-time distributed systems, and fault- and intrusion-tolerant distributed algorithms.



Jérémie Decouchant received the PhD degree in computer science from the Grenoble-Alpes University, France, in 2015. He is an assistant professor at TU Delft, The Netherlands. His research interests include resilient distributed computing, privacy-preserving systems and their application to blockchain, genomics, and machine learning.



Vincent Rahli received the PhD degree in computer science from Heriot-Watt University, U.K. He is a senior lecturer at the University of Birmingham, U.K. His research focuses on designing, formalizing, and using type theories and on the verification of distributed systems using proof assistants.



Paulo Esteves-Veríssimo (Fellow, IEEE) is a professor at the KAUST University (KSA), Saudi Arabia, and director of the Resilient Computing and Cybersecurity Center (RC3 - <https://rc3.kaust.edu.sa/>). He was a member of the Sci&Tech. Comm. of ECSO EU Cyber Security Organisation, chair of IFIP WG 10.4 on Dependable Comp. and F/T, and vice-chair of the Steer. Comm. of the DSN conference. He is fellow of ACM, and associate editor of the *IEEE Transactions on Emerging Topics in Computing* journal,

author of more than 200 peer-refereed publications and co-author of five books. He is currently interested in resilient computing and its potential to improve classic cybersecurity techniques, in SDN-based infrastructures; autonomous vehicles from earth to space; distributed control systems; digital health and genomics; or blockchain and cryptocurrencies.

▷ **For more information on this or any other computing topic, please visit our Digital Library at www.computer.org/csdl.**